

# Regression

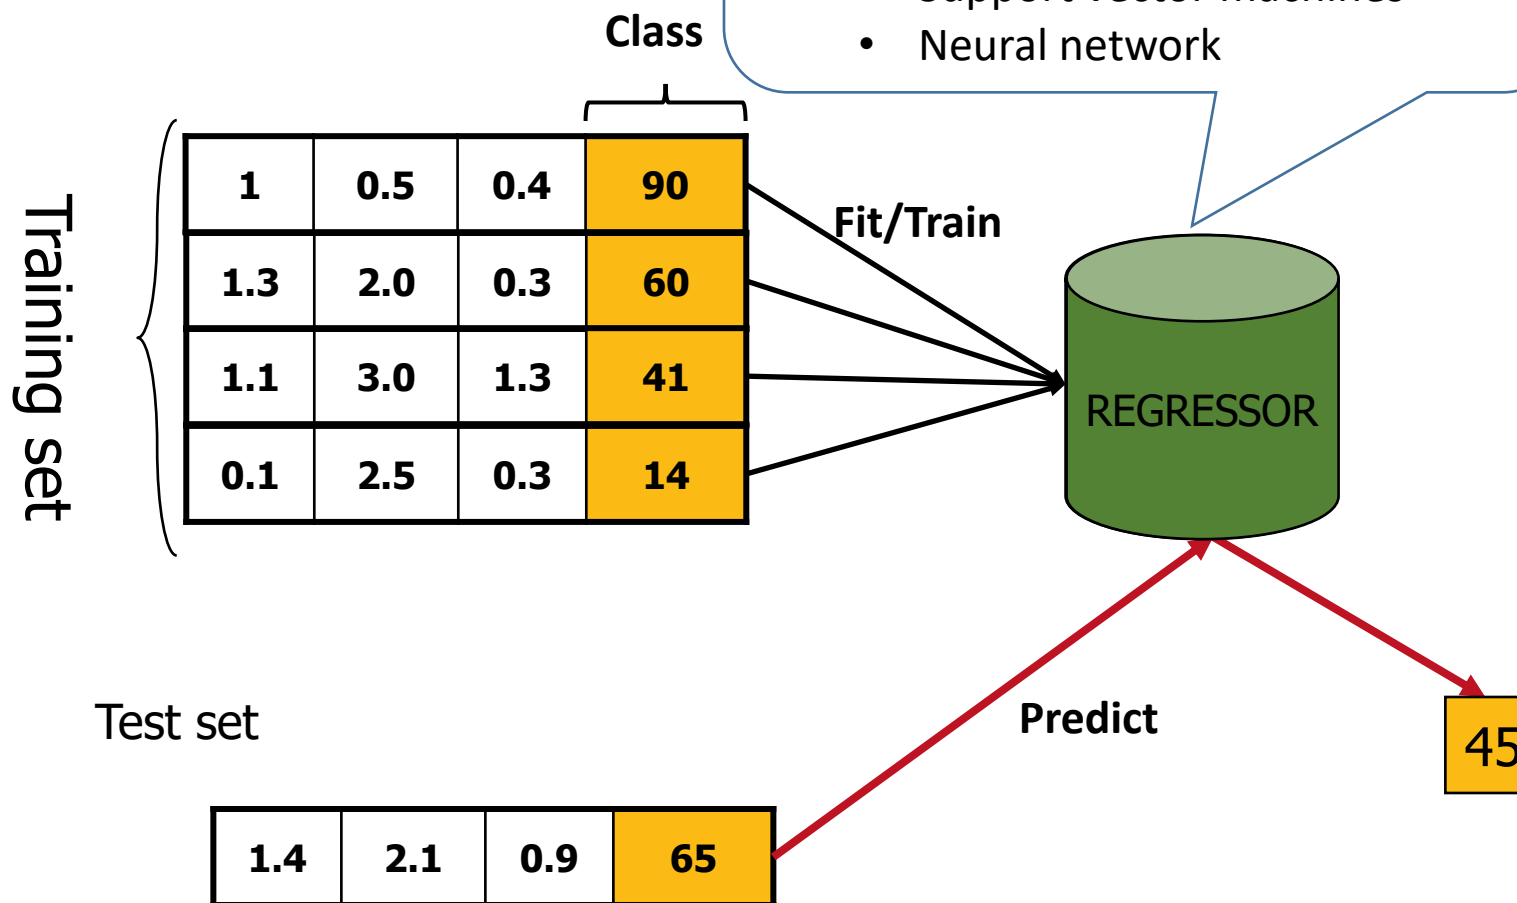
## module 12

# Regression

- Goal: Predict **numeric** target (outcome) variable
  - Examples: age, price, etc...
- Each row is a case (customer, student, applicant)
- Each column is a variable

# Regression

- Some regressors are interpretable
  - OLS regression, Lasso
  - Decision Tree Regressors
- Some aren't
  - Support vector machines
  - Neural network



# Regression for Data Exploration

Case study similar to project

# Today's data set

- <https://www.kaggle.com/uciml/adult-census-income>
- This data was extracted from the [1994 Census bureau database](#) by Ronny Kohavi and Barry Becker (Data Mining and Visualization, Silicon Graphics). *The prediction task is to determine whether a person makes over \$50K a year.*
- The cleaning is the same as in Module 9
- Today's goal: **predict age**

# The LASSO regressor

Linear prediction:  $\hat{y}_i = \mathbf{w} \cdot \mathbf{x}_i$

Minimize this function:

$$L = \sum_i \frac{1}{2} (\hat{y}_i - y_i)^2$$

Error magnitude

+

$$\frac{1}{2} \lambda |\mathbf{w}|_1$$

$\sum_i |w_i|$

Penalty for non-zero coefficients

Unlike regular regression, the LASSO regressor will try to minimize the number of predictors used.

# LASSO in scikit learn

## TRAIN

```
from sklearn import linear_model  
regLasso = linear_model.Lasso()
```

```
regLasso.fit(X,Y)
```

## INTERPRET

find the attributes with a  
non-zero coefficient

```
d={X.columns[i] : regLasso.coef_[i] for i in range(0,len(X.columns)) }
```

```
s=pd.Series(d)
```

```
s[s.abs() > 0.0001] # show the attributes whose coefficient is different from 0
```

capital.gain	0.000118
capital.loss	0.001521
hours.per.week	0.018635
marital.status_Married-civ-spouse	2.260182
relationship_Husband	2.171161
workclass_Private	-0.512760
dtype: float64	



Being married to a civilian (rather than, for instance, to one in the armed forces, or AF-spouse) is positively correlated with age

# Decision Tree Regressor

```
import sklearn.tree
dt = sklearn.tree.DecisionTreeRegressor(max_depth=2)
```


```
dt.fit(X,Y)
```

```
DecisionTreeRegressor(criterion='mse', max_depth=2, max_features=None,
                      max_leaf_nodes=None, min_impurity_split=1e-07,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                      splitter='best')
```

```
import sklearn.tree as tree
from IPython.display import Image
import pydotplus

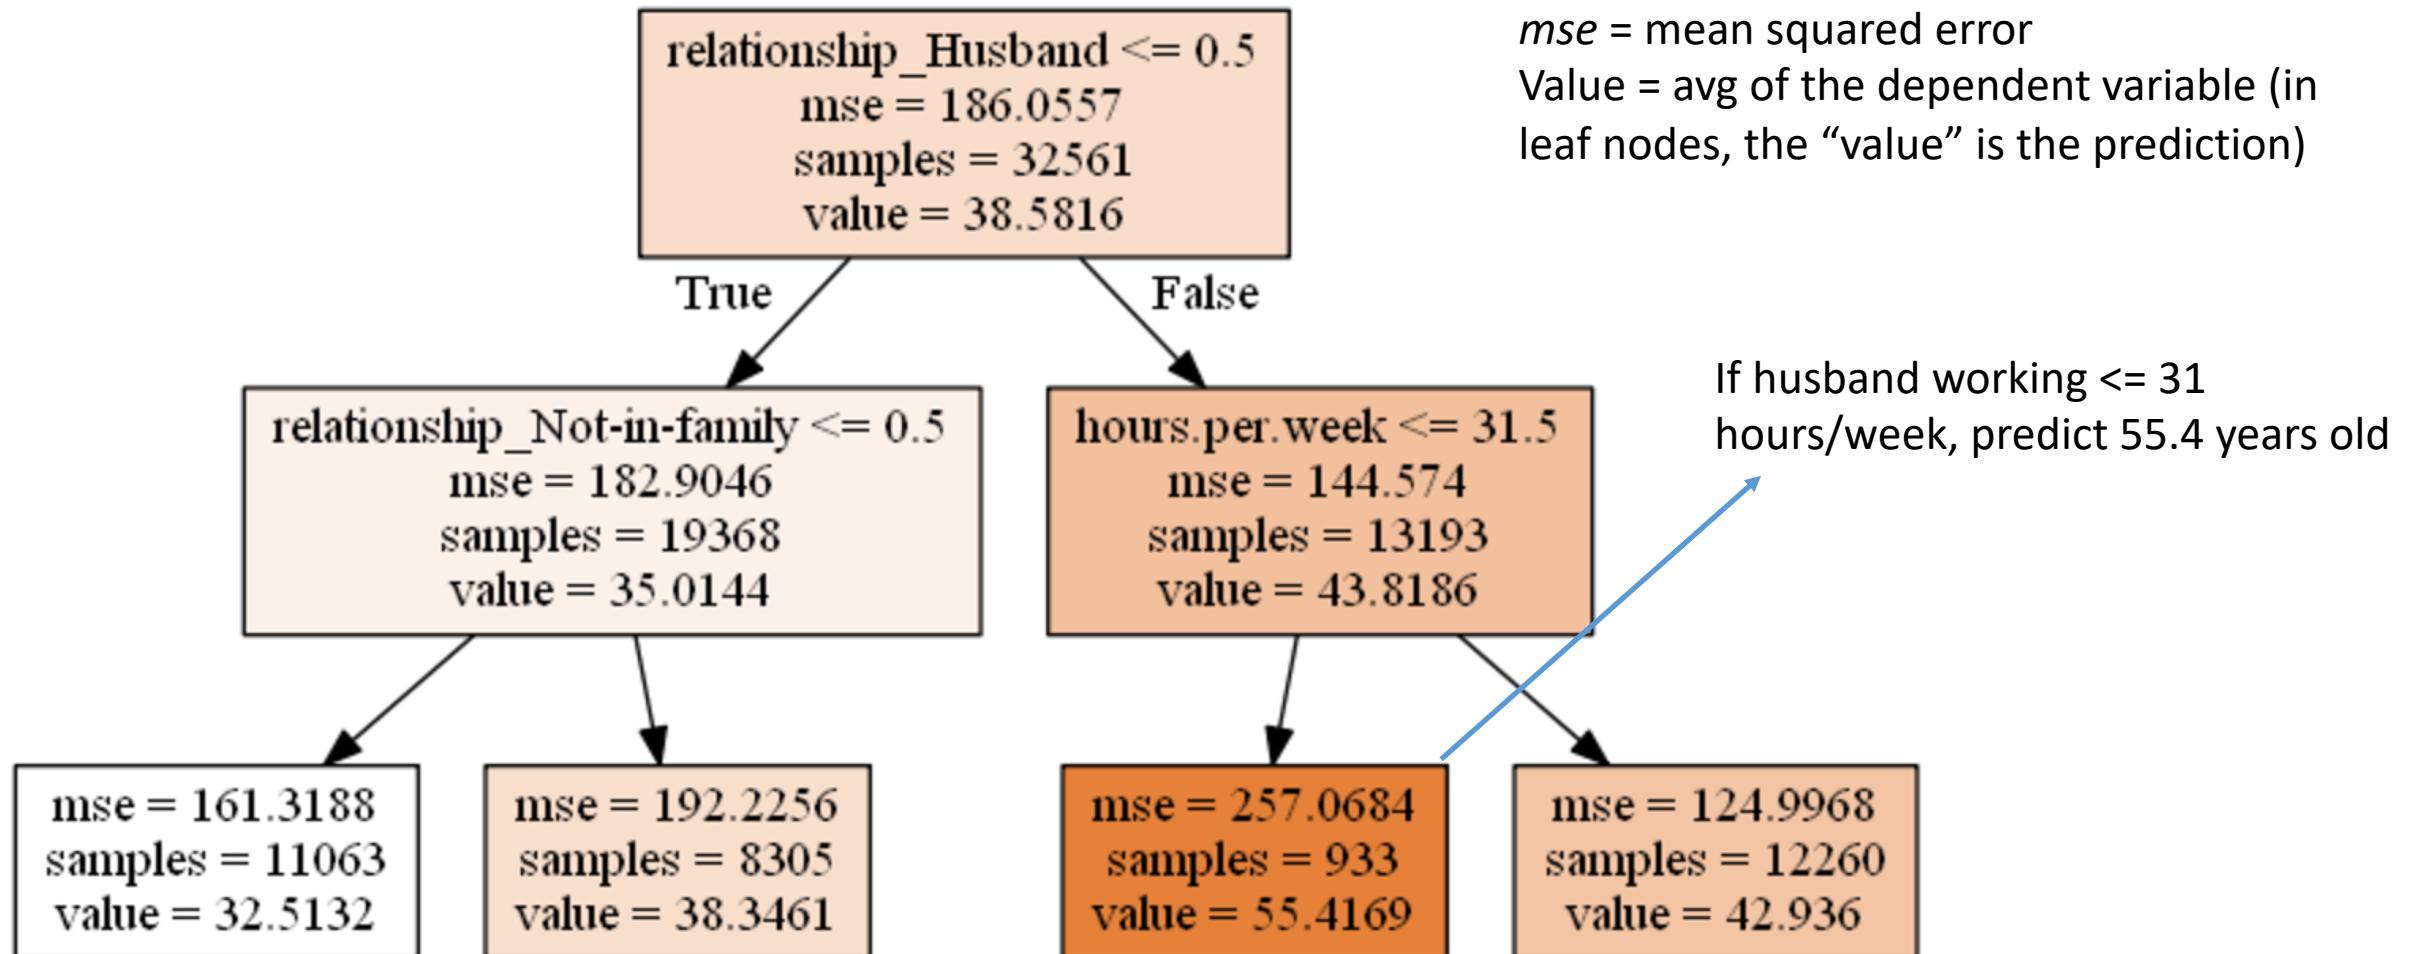
dt_feature_names = list(X.columns)
dt_target_names = np.array(Y.unique(), dtype=np.string_)
tree.export_graphviz(dt, out_file='tree.dot',
                    feature_names=dt_feature_names, class_names=dt_target_names,
                    filled=True)
graph = pydotplus.graph_from_dot_file('tree.dot')
Image(graph.create_png())
```

Same as Decision Trees  
for Classification





# Decision Tree Regressor



# Regression for Prediction

# Regression for prediction

```
from sklearn import linear_model
regLasso = linear_model.Lasso()
regLasso.fit(X_train,y_train)
y_test_pred = regLasso.predict(X_test)
```

How to measure predictive performance?

- Mean absolute deviation (MAD) =  $\frac{1}{n} \sum_{i=0}^n |y_{actual}(i) - y_{pred}(i)|$



```
(y_pred - y_test).abs().mean()
```

```
10.686668884299293
```

- Mean squared error (MSE) =  $\frac{1}{n} \sum_{i=0}^n (y_{actual}(i) - y_{pred}(i))^2$



```
((y_pred - y_test)**2).mean()
```

```
171.25305013932123
```

- ...

# Finding the best regressor

Here are a few regressors

```
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import ElasticNet
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.svm import SVR

regs = [LinearRegression(), ElasticNet(), DecisionTreeRegressor(), GradientBoostingRegressor(), MLPRegressor()]#, SVR()]
```

# Finding the best regressor

We can use Cross-validation to find the best one

```
from sklearn.model_selection import KFold

minMAD = +10000000
nfolds = 3
bestREG = ''
for reg in regs:
    kf = KFold(n_splits=nfolds, random_state=2, shuffle=True)
    thisMAD = -sklearn.model_selection.cross_val_score(reg,X,Y,cv=kf,scoring='neg_mean_absolute_error').mean()
    if thisMAD < minMAD:
        bestMAD = reg
        minMAD = thisMAD
```

Other scoring methods are reported [here](#)